

2.7 SYMBOL TABLES

- a A comment begins with `/*` and includes all characters until the end of that line
- b A comment begins with `/**` and includes all characters through the next occurrence of the character sequence `*/`

Exercise 2.6.2 Extend the lexical analyzer in Section 2.6.5 to recognize the relational operators

Exercise 2.6.3 Extend the lexical analyzer in Section 2.6.5 to recognize floating point numbers such as `2.314` and `5.`

2.7 Symbol Tables

Symbol tables are data structures that are used by compilers to hold information about source program constructs. The information is collected incrementally by the analysis phases of a compiler and used by the synthesis phases to generate the target code. Entries in the symbol table contain information about an identifier such as its character string or lexeme, its type, its position in storage, and any other relevant information. Symbol tables typically need to support multiple declarations of the same identifier within a program.

From Section 1.6.1, the scope of a declaration is the portion of a program to which the declaration applies. We shall implement scopes by setting up a separate symbol table for each scope. A program block with declarations⁸ will have its own symbol table with an entry for each declaration in the block. This approach also works for other constructs that set up scopes: for example, a class would have its own table with an entry for each field and method.

This section contains a symbol table module suitable for use with the Java translator fragments in this chapter. The module will be used as is when we put together the translator in Appendix A. Meanwhile, for simplicity, the main example of this section is a stripped-down language with just the key constructs that touch symbol tables: namely, blocks, declarations, and factors. All of the other statement and expression constructs are omitted so we can focus on the symbol table operations. A program consists of blocks with optional declarations and statements consisting of single identifiers. Each such statement represents a use of the identifier. Here is a sample program in this language:

```
int x; char y; bool y; x; y; x; y; 2.7
```

The examples of block structure in Section 1.6.3 dealt with the definitions and uses of names: the input `2.7` consists solely of definitions and uses of names.

The task we shall perform is to print a revised program in which the declarations have been removed and each statement has its identifier followed by a colon and its type.

⁸In C, for instance, program blocks are either functions or sections of functions that are separated by curly braces and that have one or more declarations within them.

Who Creates Symbol Table Entries

Symbol table entries are created and used during the analysis phase by the lexical analyzer, the parser, and the semantic analyzer. In this chapter we have the parser create entries. With its knowledge of the syntactic structure of a program, a parser is often in a better position than the lexical analyzer to distinguish among different declarations of an identifier.

In some cases, a lexical analyzer can create a symbol table entry as soon as it sees the characters that make up a lexeme. More often, the lexical analyzer can only return to the parser a token, say **id**, along with a pointer to the lexeme. Only the parser, however, can decide whether to use a previously created symbol table entry or create a new one for the identifier.

Example 2.14 On the above input, 2.7, the goal is to produce

```
x int y bool    x int y char
```

The first *x* and *y* are from the inner block of input 2.7. Since this use of *x* refers to the declaration of *x* in the outer block, it is followed by `int`, the type of that declaration. The use of *y* in the inner block refers to the declaration of *y* in that very block and therefore has boolean type. We also see the uses of *x* and *y* in the outer block, with their types, as given by declarations of the outer block: integer and character, respectively. □

2.7.1 Symbol Table Per Scope

The term *scope* of identifier *x* really refers to the scope of a particular declaration of *x*. The term *scope* by itself refers to a portion of a program that is the scope of one or more declarations.

Scopes are important, because the same identifier can be declared for different purposes in different parts of a program. Common names like *i* and *x* often have multiple uses. As another example, subclasses can redeclare a method name to override a method in a superclass.

If blocks can be nested, several declarations of the same identifier can appear within a single block. The following syntax results in nested blocks when *stmts* can generate a block:

```
block → ' ' decls stmts ' '
```

We quote curly braces in the syntax to distinguish them from curly braces for semantic actions. With the grammar in Fig. 2.38, *decls* generates an optional sequence of declarations and *stmts* generates an optional sequence of statements.

Optimization of Symbol Tables for Blocks

Implementations of symbol tables for blocks can take advantage of the most closely nested rule. Nesting ensures that the chain of applicable symbol tables forms a stack. At the top of the stack is the table for the current block. Below it in the stack are the tables for the enclosing blocks. Thus, symbol tables can be allocated and deallocated in a stack-like fashion.

Some compilers maintain a single hash table of accessible entries—that is, of entries that are not hidden by a declaration in a nested block. Such a hash table supports essentially constant time lookups at the expense of inserting and deleting entries on block entry and exit. Upon exit from a block B , the compiler must undo any changes to the hash table due to declarations in block B . It can do so by using an auxiliary stack to keep track of changes to the hash table while block B is processed.

Moreover, a statement can be a block, so our language allows nested blocks where an identifier can be redeclared.

The *most closely nested* rule for blocks is that an identifier x is in the scope of the most closely nested declaration of x —that is, the declaration of x found by examining blocks inside out, starting with the block in which x appears.

Example 2.15 The following pseudocode uses subscripts to distinguish among distinct declarations of the same identifier.

```

1  {  int  $x_1$   int  $y_1$ 
2      {  int  $w_2$   bool  $y_2$   int  $z_2$ 
3           $w_2$             $x_1$             $y_2$             $z_2$ 
4      }
5           $w_0$             $x_1$             $y_1$ 
6  }
```

The subscript is not part of an identifier; it is in fact the line number of the declaration that applies to the identifier. Thus, all occurrences of x are within the scope of the declaration on line 1. The occurrence of y on line 3 is in the scope of the declaration of y on line 2 since y is redeclared within the inner block. The occurrence of y on line 5, however, is within the scope of the declaration of y on line 1.

The occurrence of w on line 5 is presumably within the scope of a declaration of w outside this program fragment; its subscript 0 denotes a declaration that is global or external to this block.

Finally, z is declared and used within the nested block, but cannot be used on line 5, since the nested declaration applies only to the nested block. \square

CHAPTER 2 A SIMPLE SYNTAX DIRECTED TRANSLATOR

The most closely nested rule for blocks can be implemented by chaining symbol tables. That is, the table for a nested block points to the table for its enclosing block.

Example 2.16 Figure 2.36 shows symbol tables for the pseudocode in Example 2.15. B_1 is for the block starting on line 1 and B_2 is for the block starting at line 2. At the top of the figure is an additional symbol table B_0 for any global or default declarations provided by the language. During the time that we are analyzing lines 2 through 4, the environment is represented by a reference to the lowest symbol table—the one for B_2 . When we move to line 5, the symbol table for B_2 becomes inaccessible, and the environment refers instead to the symbol table for B_1 , from which we can reach the global symbol table, but not the table for B_2 . □

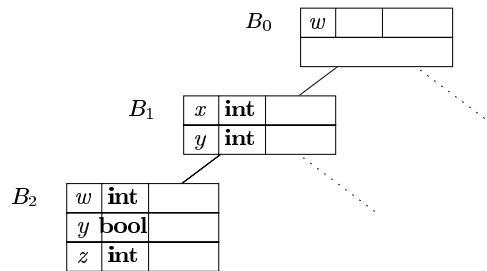


Figure 2.36 Chained symbol tables for Example 2.15

The Java implementation of chained symbol tables in Fig. 2.37 defines a class `Env`, short for *environment*.⁹ Class `Env` supports three operations:

Create a new symbol table. The constructor `Env p` on lines 6 through 8 of Fig. 2.37 creates an `Env` object with a hash table named `table`. The object is chained to the environment valued parameter `p` by setting field `prev` to `p`. Although it is the `Env` objects that form a chain, it is convenient to talk of the tables being chained.

Put a new entry in the current table. The hash table holds key-value pairs, where

The *key* is a string, or rather a reference to a string. We could alternatively use references to token objects for identifiers as keys.

The *value* is an entry of class `Symbol`. The code on lines 9 through 11 does not need to know the structure of an entry; that is, the code is independent of the fields and methods in class `Symbol`.

⁹ Environment is another term for the collection of symbol tables that are relevant at a point in the program.

2.7 SYMBOL TABLES

```
1 package symbols File Env.java
2 import java.util
3 public class Env
4     private Hashtable table
5     protected Env prev
6
7     public Env Env p
8         table = new Hashtable()
9         prev = p
10
11     public void put(String s, Symbol sym)
12         table.put(s, sym)
13
14     public Symbol get(String s)
15         for Env e = this; e != null; e = e.prev
16             Symbol found = e.table.get(s)
17             if found != null return found
18         return null
19
```

Figure 2.37 Class *Env* implements chained symbol tables

Get an entry for an identifier by searching the chain of tables starting with the table for the current block. The code for this operation on lines 12 through 18 returns either a symbol table entry or null.

Chaining of symbol tables results in a tree structure since more than one block can be nested inside an enclosing block. The dotted lines in Fig. 2.36 are a reminder that chained symbol tables can form a tree.

2.7.2 The Use of Symbol Tables

In effect, the role of a symbol table is to pass information from declarations to uses. A semantic action puts information about identifier x into the symbol table when the declaration of x is analyzed. Subsequently, a semantic action associated with a production such as *factor* **id** gets information about the identifier from the symbol table. Since the translation of an expression E_1 **op** E_2 for a typical operator **op** depends only on the translations of E_1 and E_2 and does not directly depend on the symbol table, we can add any number of operators without changing the basic flow of information from declarations to uses through the symbol table.

Example 2.17 The translation scheme in Fig. 2.38 illustrates how class *Env* can be used. The translation scheme concentrates on scopes, declarations, and

CHAPTER 2 A SIMPLE SYNTAX DIRECTED TRANSLATOR

uses It implements the translation described in Example 2.14. As noted earlier on input

<i>program</i>		{ <i>top</i> null }
	<i>block</i>	
<i>block</i>	' '	{ <i>saved</i> <i>top</i> <i>top</i> new <i>Env</i> <i>top</i> print }
	<i>decls stmts</i> ' '	{ <i>top</i> <i>saved</i> print }
<i>decls</i>	<i>decls decl</i>	
<i>decl</i>	type id	{ <i>s</i> new <i>Symbol</i> <i>s</i> <i>type</i> type <i>lexeme</i> <i>top</i> put id <i>lexeme</i> <i>s</i> }
<i>stmts</i>	<i>stmts stmt</i>	
<i>stmt</i>	<i>block</i>	
	<i>factor</i>	{ print }
<i>factor</i>	id	{ <i>s</i> <i>top</i> get id <i>lexeme</i> print id <i>lexeme</i> print } print <i>s</i> <i>type</i>

Figure 2.38 The use of symbol tables for translating a language with blocks

```
int x char y bool y x y x y
```

the translation scheme strips the declarations and produces

```
x int y bool x int y char
```

Notice that the bodies of the productions have been aligned in Fig. 2.38 so that all the grammar symbols appear in one column and all the actions in a second column. As a result, components of the body are often spread over several lines.

Now consider the semantic actions. The translation scheme creates and discards symbol tables upon block entry and exit, respectively. Variable *top* denotes the top table at the head of a chain of tables. The first production of

2.8 INTERMEDIATE CODE GENERATION

the underlying grammar is *program* \rightarrow *block*. The semantic action before *block* initializes *top* to **null** with no entries.

The second production *block* \rightarrow *' decls stmts ' ' ' has actions upon block entry and exit. On block entry, before decls, a semantic action saves a reference to the current table using a local variable saved. Each use of this production has its own local variable saved, distinct from the local variable for any other use of this production. In a recursive descent parser, saved would be local to the procedure for block. The treatment of local variables of a recursive function is discussed in Section 7.2. The code*

```
top = new Env top
```

sets variable *top* to a newly created new table that is chained to the previous value of *top* just before block entry. Variable *top* is an object of class *Env*; the code for the constructor *Env* appears in Fig. 2.37.

On block exit, after *' ' ' a semantic action restores *top* to its value saved on block entry. In effect, the tables form a stack; restoring *top* to its saved value pops the effect of the declarations in the block.¹⁰ Thus, the declarations in the block are not visible outside the block.*

A declaration *decl* \rightarrow **type** *id* results in a new entry for the declared identifier. We assume that tokens **type** and **id** each have an associated attribute which is the type and lexeme, respectively, of the declared identifier. We shall not go into all the fields of a symbol object *s*, but we assume that there is a field *type* that gives the type of the symbol. We create a new symbol object *s* and assign its type properly by *s.type = type lexeme*. The complete entry is put into the top symbol table by *top.put id lexeme s*.

The semantic action in the production *factor* \rightarrow **id** uses the symbol table to get the entry for the identifier. The *get* operation searches for the first entry in the chain of tables, starting with *top*. The retrieved entry contains any information needed about the identifier, such as the type of the identifier. \square

2.8 Intermediate Code Generation

The front end of a compiler constructs an intermediate representation of the source program from which the back end generates the target program. In this section, we consider intermediate representations for expressions and statements, and give tutorial examples of how to produce such representations.

2.8.1 Two Kinds of Intermediate Representations

As was suggested in Section 2.1 and especially Fig. 2.4, the two most important intermediate representations are

¹⁰Instead of explicitly saving and restoring tables, we could alternatively add static operations *push* and *pop* to class