

## Chapter 5

# Syntax Directed Translation

This chapter develops the theme of Section 2.3 – the translation of languages guided by context free grammars. The translation techniques in this chapter will be applied in Chapter 6 to type checking and intermediate code generation. The techniques are also useful for implementing little languages for specialized tasks: this chapter includes an example from typesetting.

We associate information with a language construct by attaching *attributes* to the grammar symbols, representing the construct, as discussed in Section 2.3.2. A syntax directed definition specifies the values of attributes by associating semantic rules with the grammar productions. For example, an infix to postfix translator might have a production and rule

PRODUCTION	SEMANTIC RULE	
$E \rightarrow E_1 T$	$E.code \leftarrow E_1.code \parallel T.code \parallel ' '$	5.1

This production has two nonterminals  $E$  and  $T$ ; the subscript in  $E_1$  distinguishes the occurrence of  $E$  in the production body from the occurrence of  $E$  as the head. Both  $E$  and  $T$  have a string valued attribute *code*. The semantic rule specifies that the string  $E.code$  is formed by concatenating  $E_1.code$ ,  $T.code$ , and the character ' '. While the rule makes it explicit that the translation of  $E$  is built up from the translations of  $E_1$ ,  $T$ , and ' ', it may be inefficient to implement the translation directly by manipulating strings.

From Section 2.3.5, a syntax directed translation scheme embeds program fragments called semantic actions within production bodies, as in

$$E \rightarrow E_1 T \{ \text{print } ' ' \} \quad 5.2$$

By convention, semantic actions are enclosed within curly braces. If curly braces occur as grammar symbols, we enclose them within single quotes, as in

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

'{' and '}' The position of a semantic action in a production body determines the order in which the action is executed. In production 5.2 the action occurs at the end, after all the grammar symbols. In general, semantic actions may occur at any position in a production body.

Between the two notations, syntax directed definitions can be more readable and hence more useful for specifications. However, translation schemes can be more efficient and hence more useful for implementations.

The most general approach to syntax directed translation is to construct a parse tree or a syntax tree and then to compute the values of attributes at the nodes of the tree by visiting the nodes of the tree. In many cases, translation can be done during parsing, without building an explicit tree. We shall therefore study a class of syntax directed translations called L attributed translations.

L for left to right, which encompass virtually all translations that can be performed during parsing. We also study a smaller class called S attributed translations. S for synthesized, which can be performed easily in connection with a bottom up parse.

### 5.1 Syntax Directed Definitions

A *syntax directed definition* (SDD) is a context free grammar together with attributes and rules. Attributes are associated with grammar symbols and rules are associated with productions. If  $X$  is a symbol and  $a$  is one of its attributes, then we write  $X.a$  to denote the value of  $a$  at a particular parse tree node labeled  $X$ . If we implement the nodes of the parse tree by records or objects, then the attributes of  $X$  can be implemented by data fields in the records that represent the nodes for  $X$ . Attributes may be of any kind: numbers, types, table references, or strings, for instance. The strings may even be long sequences of code, say code in the intermediate language used by a compiler.

#### 5.1.1 Inherited and Synthesized Attributes

We shall deal with two kinds of attributes for nonterminals.

1. A *synthesized attribute* for a nonterminal  $A$  at a parse tree node  $N$  is defined by a semantic rule associated with the production at  $N$ . Note that the production must have  $A$  as its head. A synthesized attribute at node  $N$  is defined only in terms of attribute values at the children of  $N$  and at  $N$  itself.
2. An *inherited attribute* for a nonterminal  $B$  at a parse tree node  $N$  is defined by a semantic rule associated with the production at the parent of  $N$ . Note that the production must have  $B$  as a symbol in its body. An inherited attribute at node  $N$  is defined only in terms of attribute values at  $N$ 's parent,  $N$  itself, and  $N$ 's siblings.

### An Alternative Definition of Inherited Attributes

No additional translations are enabled if we allow an inherited attribute  $B.c$  at a node  $N$  to be defined in terms of attribute values at the children of  $N$  as well as at  $N$  itself, at its parent, and at its siblings. Such rules can be simulated by creating additional attributes of  $B$ , say  $B.c_1, B.c_2$ . These are synthesized attributes that copy the needed attributes of the children of the node labeled  $B$ . We then compute  $B.c$  as an inherited attribute using the attributes  $B.c_1, B.c_2$  in place of attributes at the children. Such attributes are rarely needed in practice.

While we do not allow an inherited attribute at node  $N$  to be defined in terms of attribute values at the children of node  $N$ , we do allow a synthesized attribute at node  $N$  to be defined in terms of inherited attribute values at node  $N$  itself.

Terminals can have synthesized attributes but not inherited attributes. Attributes for terminals have lexical values that are supplied by the lexical analyzer; there are no semantic rules in the SDD itself for computing the value of an attribute for a terminal.

**Example 5.1** The SDD in Fig. 5.1 is based on our familiar grammar for arithmetic expressions with operators and. It evaluates expressions terminated by an endmarker **n**. In the SDD, each of the nonterminals has a single synthesized attribute called *val*. We also suppose that the terminal **digit** has a synthesized attribute *lexval*, which is an integer value returned by the lexical analyzer.

	PRODUCTION	SEMANTIC RULES
1	$L \rightarrow E \mathbf{n}$	$L.val \leftarrow E.val$
2	$E \rightarrow E_1 T$	$E.val \leftarrow E_1.val + T.val$
3	$E \rightarrow T$	$E.val \leftarrow T.val$
4	$T \rightarrow T_1 F$	$T.val \leftarrow T_1.val * F.val$
5	$T \rightarrow F$	$T.val \leftarrow F.val$
6	$F \rightarrow E$	$F.val \leftarrow E.val$
7	$F \rightarrow \mathbf{digit}$	$F.val \leftarrow \mathbf{digit.lexval}$

Figure 5.1 Syntax directed definition of a simple desk calculator

The rule for production 1,  $L \rightarrow E \mathbf{n}$ , sets  $L.val$  to  $E.val$ , which we shall see is the numerical value of the entire expression.

Production 2,  $E \rightarrow E_1 T$ , also has one rule, which computes the *val* attribute for the head  $E$  as the sum of the values at  $E_1$  and  $T$ . At any parse

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

tree node  $N$  labeled  $E$  the value of  $val$  for  $E$  is the sum of the values of  $val$  at the children of node  $N$  labeled  $E$  and  $T$

Production 3  $E \rightarrow T$  has a single rule that defines the value of  $val$  for  $E$  to be the same as the value of  $val$  at the child for  $T$ . Production 4 is similar to the second production: its rule multiplies the values at the children instead of adding them. The rules for productions 5 and 6 copy values at a child like that for the third production. Production 7 gives  $F.val$  the value of a digit—that is, the numerical value of the token **digit** that the lexical analyzer returned.  $\square$

An SDD that involves only synthesized attributes is called *S attributed*; the SDD in Fig. 5.1 has this property. In an S attributed SDD, each rule computes an attribute for the nonterminal at the head of a production from attributes taken from the body of the production.

For simplicity, the examples in this section have semantic rules without side effects. In practice, it is convenient to allow SDDs to have limited side effects, such as printing the result computed by a desk calculator or interacting with a symbol table. Once the order of evaluation of attributes is discussed in Section 5.2, we shall allow semantic rules to compute arbitrary functions, possibly involving side effects.

An S attributed SDD can be implemented naturally in conjunction with an LR parser. In fact, the SDD in Fig. 5.1 mirrors the Yacc program of Fig. 4.58, which illustrates translation during LR parsing. The difference is that, in the rule for production 1, the Yacc program prints the value  $E.val$  as a side effect instead of defining the attribute  $L.val$ .

An SDD without side effects is sometimes called an *attribute grammar*. The rules in an attribute grammar define the value of an attribute purely in terms of the values of other attributes and constants.

### 5.1.2 Evaluating an SDD at the Nodes of a Parse Tree

To visualize the translation specified by an SDD, it helps to work with parse trees, even though a translator need not actually build a parse tree. Imagine therefore that the rules of an SDD are applied by first constructing a parse tree and then using the rules to evaluate all of the attributes at each of the nodes of the parse tree. A parse tree showing the values of its attributes is called an *annotated parse tree*.

How do we construct an annotated parse tree? In what order do we evaluate attributes? Before we can evaluate an attribute at a node of a parse tree, we must evaluate all the attributes upon which its value depends. For example, if all attributes are synthesized—as in Example 5.1—then we must evaluate the  $val$  attributes at all of the children of a node before we can evaluate the  $val$  attribute at the node itself.

With synthesized attributes, we can evaluate attributes in any bottom-up order, such as that of a postorder traversal of the parse tree; the evaluation of S attributed definitions is discussed in Section 5.2.3.

## 5.1 SYNTAX DIRECTED DEFINITIONS

For SDDs with both inherited and synthesized attributes there is no guarantee that there is even one order in which to evaluate attributes at nodes. For instance, consider nonterminals  $A$  and  $B$  with synthesized and inherited attributes  $A.s$  and  $B.i$  respectively, along with the production and rules

PRODUCTION	SEMANTIC RULES
$A \rightarrow B$	$A.s = B.i$
	$B.i = A.s + 1$

These rules are circular: it is impossible to evaluate either  $A.s$  at a node  $N$  or  $B.i$  at the child of  $N$  without first evaluating the other. The circular dependency of  $A.s$  and  $B.i$  at some pair of nodes in a parse tree is suggested by Fig. 5.2




Figure 5.2 The circular dependency of  $A.s$  and  $B.i$  on one another

It is computationally difficult to determine whether or not there exist any circularities in any of the parse trees that a given SDD could have to translate.<sup>1</sup> Fortunately, there are useful subclasses of SDDs that are sufficient to guarantee that an order of evaluation exists, as we shall see in Section 5.2.

**Example 5.2** Figure 5.3 shows an annotated parse tree for the input string 3 5 4 n constructed using the grammar and rules of Fig. 5.1. The values of *lexval* are presumed supplied by the lexical analyzer. Each of the nodes for the nonterminals has attribute *val* computed in a bottom-up order, and we see the resulting values associated with each node. For instance, at the node with a child labeled *a*, after computing  $T.val = 3$  and  $F.val = 5$  at its first and third children, we apply the rule that says  $T.val$  is the product of these two values or 15. □

Inherited attributes are useful when the structure of a parse tree does not match the abstract syntax of the source code. The next example shows how inherited attributes can be used to overcome such a mismatch due to a grammar designed for parsing rather than translation.

<sup>1</sup>Without going into details, while the problem is decidable, it cannot be solved by a polynomial time algorithm, even if  $\mathcal{P} = \mathcal{NP}$ , since it has exponential time complexity.

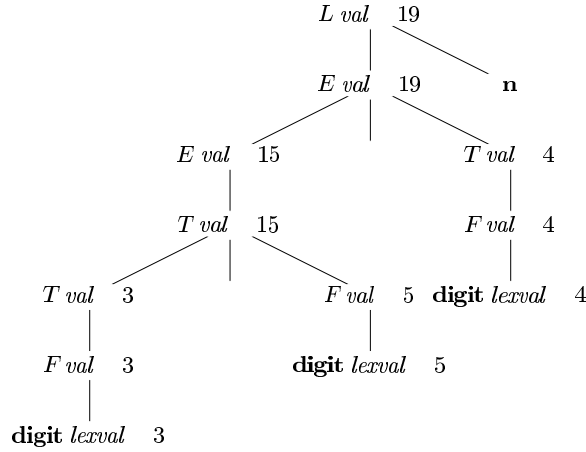


Figure 5.3 Annotated parse tree for 3 \* 5 + 4 \* n

**Example 5.3** The SDD in Fig. 5.4 computes terms like 3 \* 5 and 3 \* 5 + 7. The top down parse of input 3 \* 5 begins with the production  $T \rightarrow FT'$ . Here  $F$  generates the digit 3 but the operator \* is generated by  $T'$ . Thus the left operand 3 appears in a different subtree of the parse tree from \*. An inherited attribute will therefore be used to pass the operand to the operator.

The grammar in this example is an excerpt from a non left recursive version of the familiar expression grammar we used such a grammar as a running example to illustrate top down parsing in Section 4.4.

	PRODUCTION	SEMANTIC RULES
1	$T \rightarrow FT'$	$T'_{inh} = F_{val}$ $T_{val} = T'_{syn}$
2	$T' \rightarrow FT'_1$	$T'_1_{inh} = T'_{inh} = F_{val}$ $T'_{syn} = T'_1_{syn}$
3	$T' \rightarrow T'$	$T'_{syn} = T'_{inh}$
4	$F \rightarrow \text{digit}$	$F_{val} = \text{digit}_{lexval}$

Figure 5.4 An SDD based on a grammar suitable for top down parsing

Each of the nonterminals  $T$  and  $F$  has a synthesized attribute *val* the terminal **digit** has a synthesized attribute *lexval*. The nonterminal  $T'$  has two attributes an inherited attribute *inh* and a synthesized attribute *syn*.

## 5.1 SYNTAX DIRECTED DEFINITIONS

The semantic rules are based on the idea that the left operand of the operator is inherited. More precisely, the head  $T'$  of the production  $T' \rightarrow F T'_1$  inherits the left operand of  $\rightarrow$  in the production body. Given a term  $x \rightarrow y \rightarrow z$ , the root of the subtree for  $y \rightarrow z$  inherits  $x$ . Then, the root of the subtree for  $z$  inherits the value of  $x \rightarrow y$  and so on, if there are more factors in the term. Once all the factors have been accumulated, the result is passed back up the tree using synthesized attributes.

To see how the semantic rules are used, consider the annotated parse tree for  $3 \rightarrow 5$  in Fig. 5.5. The leftmost leaf in the parse tree, labeled **digit**, has attribute value  $lexval = 3$ , where the 3 is supplied by the lexical analyzer. Its parent is for production 4,  $F \rightarrow \text{digit}$ . The only semantic rule associated with this production defines  $F.val = \text{digit}.lexval$ , which equals 3.




Figure 5.5 Annotated parse tree for  $3 \rightarrow 5$

At the second child of the root, the inherited attribute  $T'.inh$  is defined by the semantic rule  $T'.inh = F.val$  associated with production 1. Thus, the left operand 3 for the  $\rightarrow$  operator is passed from left to right across the children of the root.

The production at the node for  $T'$  is  $T' \rightarrow F T'_1$ . We retain the subscript 1 in the annotated parse tree to distinguish between the two nodes for  $T'$ . The inherited attribute  $T'_1.inh$  is defined by the semantic rule  $T'_1.inh = T'.inh \rightarrow F.val$  associated with production 2.

With  $T'.inh = 3$  and  $F.val = 5$ , we get  $T'_1.inh = 15$ . At the lower node for  $T'_1$ , the production is  $T' \rightarrow F T'_1$ . The semantic rule  $T'.syn = T'.inh$  defines  $T'_1.syn = 15$ . The  $syn$  attributes at the nodes for  $T'$  pass the value 15 up the tree to the node for  $T$ , where  $T.val = 15$ .  $\square$

### 5.1.3 Exercises for Section 5.1

**Exercise 5.1.1** For the SDD of Fig. 5.1, give annotated parse trees for the following expressions:

- a.  $3 \rightarrow 4 \rightarrow 5 \rightarrow 6 \rightarrow n$

```

b 1 2 3 4 5 n
c 9 8 7 6 5 4n

```

**Exercise 5.1.2** Extend the SDD of Fig. 5.4 to handle expressions as in Fig. 5.1

**Exercise 5.1.3** Repeat Exercise 5.1.1 using your SDD from Exercise 5.1.2

## 5.2 Evaluation Orders for SDDs

Dependency graphs are a useful tool for determining an evaluation order for the attribute instances in a given parse tree. While an annotated parse tree shows the values of attributes, a dependency graph helps us determine how those values can be computed.

In this section, in addition to dependency graphs, we define two important classes of SDDs: the S-attributed and the more general L-attributed SDDs. The translations specified by these two classes fit well with the parsing methods we have studied, and most translations encountered in practice can be written to conform to the requirements of at least one of these classes.

### 5.2.1 Dependency Graphs

A *dependency graph* depicts the flow of information among the attribute instances in a particular parse tree: an edge from one attribute instance to another means that the value of the first is needed to compute the second. Edges express constraints implied by the semantic rules. In more detail:

For each parse tree node, say a node labeled by grammar symbol  $X$ , the dependency graph has a node for each attribute associated with  $X$ .

Suppose that a semantic rule associated with a production  $p$  defines the value of synthesized attribute  $A.b$  in terms of the value of  $X.c$ ; the rule may define  $A.b$  in terms of other attributes in addition to  $X.c$ . Then the dependency graph has an edge from  $X.c$  to  $A.b$ . More precisely, at every node  $N$  labeled  $A$  where production  $p$  is applied, create an edge to attribute  $b$  at  $N$  from the attribute  $c$  at the child of  $N$  corresponding to this instance of the symbol  $X$  in the body of the production.<sup>2</sup>

Suppose that a semantic rule associated with a production  $p$  defines the value of inherited attribute  $B.c$  in terms of the value of  $X.a$ . Then the dependency graph has an edge from  $X.a$  to  $B.c$ . For each node  $N$  labeled  $B$  that corresponds to an occurrence of this  $B$  in the body of production  $p$ , create an edge to attribute  $c$  at  $N$  from the attribute  $a$  at the node  $M$

<sup>2</sup>Since a node can have several children labeled  $X$ , we again assume that subscripts distinguish among uses of the same symbol at different places in the production.



## 5.2 EVALUATION ORDERS FOR SDD'S

that corresponds to this occurrence of  $X$ . Note that  $M$  could be either the parent or a sibling of  $N$ .

**Example 5.4** Consider the following production and rule

PRODUCTION	SEMANTIC RULE
$E \rightarrow E_1 \ T$	$E.val \leftarrow E_1.val \ T.val$

At every node  $N$  labeled  $E$  with children corresponding to the body of this production, the synthesized attribute  $val$  at  $N$  is computed using the values of  $val$  at the two children labeled  $E_1$  and  $T$ . Thus, a portion of the dependency graph for every parse tree in which this production is used looks like Fig. 5.6. As a convention, we shall show the parse tree edges as dotted lines, while the edges of the dependency graph are solid.  $\square$




Figure 5.6  $E.val$  is synthesized from  $E_1.val$  and  $T.val$

**Example 5.5** An example of a complete dependency graph appears in Fig. 5.7. The nodes of the dependency graph, represented by the numbers 1 through 9, correspond to the attributes in the annotated parse tree in Fig. 5.5.




Figure 5.7 Dependency graph for the annotated parse tree of Fig. 5.5

Nodes 1 and 2 represent the attribute  $lexval$  associated with the two leaves labeled **digit**. Nodes 3 and 4 represent the attribute  $val$  associated with the two nodes labeled  $F$ . The edges to node 3 from 1 and to node 4 from 2 result

from the semantic rule that defines  $Fval$  in terms of **digit**  $lexval$ . In fact  $Fval$  equals **digit**  $lexval$  but the edge represents dependence not equality.

Nodes 5 and 6 represent the inherited attribute  $T' inh$  associated with each of the occurrences of nonterminal  $T'$ . The edge to 5 from 3 is due to the rule  $T' inh = Fval$  which defines  $T' inh$  at the right child of the root from  $Fval$  at the left child. We see edges to 6 from node 5 for  $T' inh$  and from node 4 for  $Fval$  because these values are multiplied to evaluate the attribute  $inh$  at node 6.

Nodes 7 and 8 represent the synthesized attribute  $syn$  associated with the occurrences of  $T'$ . The edge to node 7 from 6 is due to the semantic rule  $T' syn = T' inh$  associated with production 3 in Fig. 5.4. The edge to node 8 from 7 is due to a semantic rule associated with production 2.

Finally, node 9 represents the attribute  $Tval$ . The edge to 9 from 8 is due to the semantic rule  $Tval = T' syn$  associated with production 1.  $\square$

### 5.2.2 Ordering the Evaluation of Attributes

The dependency graph characterizes the possible orders in which we can evaluate the attributes at the various nodes of a parse tree. If the dependency graph has an edge from node  $M$  to node  $N$  then the attribute corresponding to  $M$  must be evaluated before the attribute of  $N$ . Thus, the only allowable orders of evaluation are those sequences of nodes  $N_1, N_2, \dots, N_k$  such that if there is an edge of the dependency graph from  $N_i$  to  $N_j$  then  $i < j$ . Such an ordering embeds a directed graph into a linear order and is called a *topological sort* of the graph.

If there is any cycle in the graph, then there are no topological sorts; that is, there is no way to evaluate the SDD on this parse tree. If there are no cycles, however, then there is always at least one topological sort. To see why, since there are no cycles, we can surely find a node with no edge entering. For if there were no such node, we could proceed from predecessor to predecessor until we came back to some node we had already seen, yielding a cycle. Make this node the first in the topological order, remove it from the dependency graph, and repeat the process on the remaining nodes.

**Example 5.6** The dependency graph of Fig. 5.7 has no cycles. One topological sort is the order in which the nodes have already been numbered: 1, 2, ..., 9. Notice that every edge of the graph goes from a node to a higher-numbered node, so this order is surely a topological sort. There are other topological sorts as well, such as 1, 3, 5, 2, 4, 6, 7, 8, 9.  $\square$

### 5.2.3 S Attributed Definitions

As mentioned earlier, given an SDD, it is very hard to tell whether there exist any parse trees whose dependency graphs have cycles. In practice, translations can be implemented using classes of SDDs that guarantee an evaluation order.

## 5.2 EVALUATION ORDERS FOR SDD'S

since they do not permit dependency graphs with cycles. Moreover, the two classes introduced in this section can be implemented efficiently in connection with top down or bottom up parsing.

The first class is defined as follows:

An SDD is *S attributed* if every attribute is synthesized.

**Example 5.7** The SDD of Fig. 5.1 is an example of an S attributed definition. Each attribute  $L.val$ ,  $E.val$ ,  $T.val$ , and  $F.val$  is synthesized.  $\square$

When an SDD is S attributed, we can evaluate its attributes in any bottom up order of the nodes of the parse tree. It is often especially simple to evaluate the attributes by performing a postorder traversal of the parse tree and evaluating the attributes at a node  $N$  when the traversal leaves  $N$  for the last time. That is, we apply the function *postorder* defined below to the root of the parse tree (see also the box "Preorder and Postorder Traversals" in Section 2.3.4).

```

postorder  $N$  {
    for each child  $C$  of  $N$  from the left postorder  $C$ 
    evaluate the attributes associated with node  $N$ 
}
```

S attributed definitions can be implemented during bottom up parsing, since a bottom up parse corresponds to a postorder traversal. Specifically, postorder corresponds exactly to the order in which an LR parser reduces a production body to its head. This fact will be used in Section 5.4.2 to evaluate synthesized attributes and store them on the stack during LR parsing, without creating the tree nodes explicitly.

### 5.2.4 L Attributed Definitions

The second class of SDD's is called *L attributed definitions*. The idea behind this class is that, between the attributes associated with a production body, dependency graph edges can go from left to right, but not from right to left; hence, L attributed. More precisely, each attribute must be either

1. Synthesized, or
2. Inherited, but with the rules limited as follows. Suppose that there is a production  $A \rightarrow X_1X_2 \dots X_n$  and that there is an inherited attribute  $X_i.a$  computed by a rule associated with this production. Then the rule may use only
  - a. Inherited attributes associated with the head  $A$ .
  - b. Either inherited or synthesized attributes associated with the occurrences of symbols  $X_1, X_2, \dots, X_{i-1}$  located to the left of  $X_i$ .

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

- c Inherited or synthesized attributes associated with this occurrence of  $X_i$  itself but only in such a way that there are no cycles in a dependency graph formed by the attributes of this  $X_i$

**Example 5.8** The SDD in Fig. 5.4 is L attributed. To see why, consider the semantic rules for inherited attributes, which are repeated here for convenience.

PRODUCTION	SEMANTIC RULE
$T \rightarrow F T'$	$T' \text{ inh} \leftarrow F \text{ val}$
$T' \rightarrow F T'_1$	$T'_1 \text{ inh} \leftarrow T' \text{ inh}, F \text{ val}$

The first of these rules defines the inherited attribute  $T' \text{ inh}$  using only  $F \text{ val}$  and  $F$  appears to the left of  $T'$  in the production body, as required. The second rule defines  $T'_1 \text{ inh}$  using the inherited attribute  $T' \text{ inh}$  associated with the head and  $F \text{ val}$ , where  $F$  appears to the left of  $T'_1$  in the production body.

In each of these cases, the rules use information from above or from the left, as required by the class. The remaining attributes are synthesized. Hence the SDD is L attributed.  $\square$

**Example 5.9** Any SDD containing the following production and rules cannot be L attributed.

PRODUCTION	SEMANTIC RULES
$A \rightarrow B C$	$A \text{ s} \leftarrow B \text{ b}$
	$B \text{ i} \leftarrow f(C \text{ c}, A \text{ s})$

The first rule  $A \text{ s} \leftarrow B \text{ b}$  is a legitimate rule in either an S attributed or L attributed SDD. It defines a synthesized attribute  $A \text{ s}$  in terms of an attribute at a child, that is, a symbol within the production body.

The second rule defines an inherited attribute  $B \text{ i}$ , so the entire SDD cannot be S attributed. Further, although the rule is legal, the SDD cannot be L attributed, because the attribute  $C \text{ c}$  is used to help define  $B \text{ i}$ , and  $C$  is to the right of  $B$  in the production body. While attributes at siblings in a parse tree may be used in L attributed SDDs, they must be to the left of the symbol whose attribute is being defined.  $\square$

### 5.2.5 Semantic Rules with Controlled Side Effects

In practice, translations involve side effects: a desk calculator might print a result, a code generator might enter the type of an identifier into a symbol table. With SDDs, we strike a balance between attribute grammars and translation schemes. Attribute grammars have no side effects and allow any evaluation order consistent with the dependency graph. Translation schemes impose left-to-right evaluation and allow semantic actions to contain any program fragment. Translation schemes are discussed in Section 5.4.

We shall control side effects in SDDs in one of the following ways:

## 5.2 EVALUATION ORDERS FOR SDD'S

Permit incidental side effects that do not constrain attribute evaluation. In other words, permit side effects when attribute evaluation based on any topological sort of the dependency graph produces a correct translation where correct depends on the application.

Constrain the allowable evaluation orders so that the same translation is produced for any allowable order. The constraints can be thought of as implicit edges added to the dependency graph.

As an example of an incidental side effect, let us modify the desk calculator of Example 5.1 to print a result. Instead of the rule  $L \rightarrow val \leftarrow E \rightarrow val$  which saves the result in the synthesized attribute  $L \rightarrow val$ , consider

	PRODUCTION	SEMANTIC RULE
1	$L \rightarrow E \mathbf{n}$	$\text{print } E \rightarrow val$

Semantic rules that are executed for their side effects, such as  $\text{print } E \rightarrow val$ , will be treated as the definitions of dummy synthesized attributes associated with the head of the production. The modified SDD produces the same translation under any topological sort, since the print statement is executed at the end, after the result is computed into  $E \rightarrow val$ .

**Example 5.10** The SDD in Fig. 5.8 takes a simple declaration  $D$  consisting of a basic type  $T$  followed by a list  $L$  of identifiers.  $T$  can be **int** or **oat**. For each identifier on the list, the type is entered into the symbol table entry for the identifier. We assume that entering the type for one identifier does not affect the symbol table entry for any other identifier. Thus, entries can be updated in any order. This SDD does not check whether an identifier is declared more than once; it can be modified to do so.

	PRODUCTION	SEMANTIC RULES
1	$D \rightarrow T L$	$L \rightarrow inh \leftarrow T \rightarrow type$
2	$T \rightarrow \mathbf{int}$	$T \rightarrow type \leftarrow \text{integer}$
3	$T \rightarrow \mathbf{oat}$	$T \rightarrow type \leftarrow \text{oat}$
4	$L \rightarrow L_1 \mathbf{id}$	$L_1 \rightarrow inh \leftarrow L \rightarrow inh$ $\text{addType } \mathbf{id} \text{ entry } L \rightarrow inh$
5	$L \rightarrow \mathbf{id}$	$\text{addType } \mathbf{id} \text{ entry } L \rightarrow inh$

Figure 5.8 Syntax directed definition for simple type declarations

Nonterminal  $D$  represents a declaration, which, from production 1, consists of a type  $T$  followed by a list  $L$  of identifiers.  $T$  has one attribute,  $T \rightarrow type$ , which is the type in the declaration  $D$ . Nonterminal  $L$  also has one attribute, which we call  $inh$  to emphasize that it is an inherited attribute. The purpose of  $L \rightarrow inh$

is to pass the declared type down the list of identifiers so that it can be added to the appropriate symbol table entries

Productions 2 and 3 each evaluate the synthesized attribute  $T\ type$  giving it the appropriate value integer or oat This type is passed to the attribute  $L\ inh$  in the rule for production 1 Production 4 passes  $L\ inh$  down the parse tree That is the value  $L_1\ inh$  is computed at a parse tree node by copying the value of  $L\ inh$  from the parent of that node the parent corresponds to the head of the production

Productions 4 and 5 also have a rule in which a function *addType* is called with two arguments

- 1 **id entry** a lexical value that points to a symbol table object and
- 2 *L inh* the type being assigned to every identifier on the list

We suppose that function *addType* properly installs the type *L inh* as the type of the represented identifier

A dependency graph for the input string `oat id1 id2 id3` appears in Fig. 5.9. Numbers 1 through 10 represent the nodes of the dependency graph. Nodes 1, 2, and 3 represent the attribute *entry* associated with each of the leaves labeled **id**. Nodes 6, 8, and 10 are the dummy attributes that represent the application of the function *addType* to a type and one of these *entry* values.

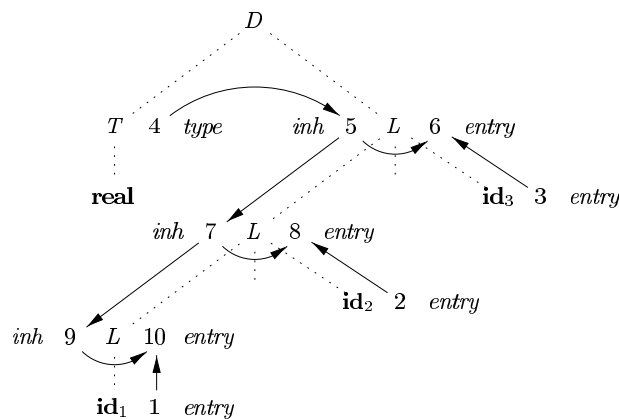


Figure 5.9 Dependency graph for a declaration `oat id1 id2 id3`

Node 4 represents the attribute *T type* and is actually where attribute evaluation begins. This type is then passed to nodes 5–7 and 9 representing *L inh* associated with each of the occurrences of the nonterminal *L*.  $\square$

## 5.2 EVALUATION ORDERS FOR SDD'S

### 5.2.6 Exercises for Section 5.2

**Exercise 5.2.1** What are all the topological sorts for the dependency graph of Fig. 5.7?

**Exercise 5.2.2** For the SDD of Fig. 5.8, give annotated parse trees for the following expressions:

```
a int a b c
b float w x y z
```

**Exercise 5.2.3** Suppose that we have a production  $A \rightarrow BCD$ . Each of the four nonterminals  $A$ ,  $B$ ,  $C$ , and  $D$  has two attributes:  $s$  is a synthesized attribute and  $i$  is an inherited attribute. For each of the sets of rules below, tell whether  $i$  the rules are consistent with an S attributed definition,  $ii$  the rules are consistent with an L attributed definition, and  $iii$  whether the rules are consistent with any evaluation order at all.

- a  $A.s \rightarrow B.i \ C.s$
- b  $A.s \rightarrow B.i \ C.s$  and  $D.i \rightarrow A.i \ B.s$
- c  $A.s \rightarrow B.s \ D.s$
- d  $A.s \rightarrow D.i \ B.i \ A.s \ C.s$ ,  $C.i \rightarrow B.s$  and  $D.i \rightarrow B.i \ C.i$

**Exercise 5.2.4** This grammar generates binary numbers with a decimal point:

$$\begin{array}{lcl} S & \rightarrow & L \mid L \mid L \\ L & \rightarrow & L \mid B \mid B \\ B & \rightarrow & 0 \mid 1 \end{array}$$

Design an L attributed SDD to compute  $S.val$ , the decimal number value of an input string. For example, the translation of string 101.101 should be the decimal number 5.625. *Hint:* use an inherited attribute  $L.side$  that tells which side of the decimal point a bit is on.

**Exercise 5.2.5** Design an S attributed SDD for the grammar and translation described in Exercise 5.2.4.

**Exercise 5.2.6** Implement Algorithm 3.23, which converts a regular expression into a nondeterministic finite automaton, by an L attributed SDD on a top-down parsable grammar. Assume that there is a token **char** representing any character, and that **char.lexval** is the character it represents. You may also assume the existence of a function *new* that returns a new state, that is, a state never before returned by this function. Use any convenient notation to specify the transitions of the NFA.

### 5.3 Applications of Syntax Directed Translation

The syntax directed translation techniques in this chapter will be applied in Chapter 6 to type checking and intermediate code generation. Here we consider selected examples to illustrate some representative SDDs.

The main application in this section is the construction of syntax trees. Since some compilers use syntax trees as an intermediate representation, a common form of SDD turns its input string into a tree. To complete the translation to intermediate code, the compiler may then walk the syntax tree using another set of rules that are in effect an SDD on the syntax tree rather than the parse tree. Chapter 6 also discusses approaches to intermediate code generation that apply an SDD without ever constructing a tree explicitly.

We consider two SDDs for constructing syntax trees for expressions. The first, an S-attributed definition, is suitable for use during bottom-up parsing. The second, an L-attributed definition, is suitable for use during top-down parsing.

The final example of this section is an L-attributed definition that deals with basic and array types.

#### 5.3.1 Construction of Syntax Trees

As discussed in Section 2.8.2, each node in a syntax tree represents a construct; the children of the node represent the meaningful components of the construct. A syntax tree node representing an expression  $E_1 \ E_2$  has label `op` and two children representing the subexpressions  $E_1$  and  $E_2$ .

We shall implement the nodes of a syntax tree by objects with a suitable number of fields. Each object will have an *op* field that is the label of the node. The objects will have additional fields as follows:

If the node is a leaf, an additional field holds the lexical value for the leaf. A constructor function *Leaf op val* creates a leaf object. Alternatively, if nodes are viewed as records, then *Leaf* returns a pointer to a new record for a leaf.

If the node is an interior node, there are as many additional fields as the node has children in the syntax tree. A constructor function *Node op c<sub>1</sub> c<sub>2</sub> ... c<sub>k</sub>* creates an object with first field *op* and *k* additional fields for the *k* children *c<sub>1</sub>* ... *c<sub>k</sub>*.

**Example 5.11** The S-attributed definition in Fig. 5.10 constructs syntax trees for a simple expression grammar involving only the binary operators `and` and `or`. As usual, these operators are at the same precedence level and are jointly left associative. All nonterminals have one synthesized attribute *node*, which represents a node of the syntax tree.

Every time the first production  $E \rightarrow E_1 \ T$  is used, its rule creates a node with `' '` for *op* and two children  $E_1$  *node* and *T node* for the subexpressions. The second production has a similar rule.



### 5.3 APPLICATIONS OF SYNTAX DIRECTED TRANSLATION

PRODUCTION				SEMANTIC RULES	
1	$E$	$E_1$	$T$	$E\ node$	<b>new</b> $Node\ ' \ ' \ E_1\ node\ T\ node$
2	$E$	$E_1$	$T$	$E\ node$	<b>new</b> $Node\ ' \ ' \ E_1\ node\ T\ node$
3	$E$	$T$		$E\ node$	$T\ node$
4	$T$	$E$		$T\ node$	$E\ node$
5	$T$	<b>id</b>		$T\ node$	<b>new</b> $Leaf\ id\ id\ entry$
6	$T$	<b>num</b>		$T\ node$	<b>new</b> $Leaf\ num\ num\ val$

Figure 5.10 Constructing syntax trees for simple expressions

For production 3  $E \rightarrow T$  no node is created since  $E\ node$  is the same as  $T\ node$ . Similarly no node is created for production 4  $T \rightarrow E$ . The value of  $T\ node$  is the same as  $E\ node$  since parentheses are used only for grouping they influence the structure of the parse tree and the syntax tree but once their job is done there is no further need to retain them in the syntax tree.

The last two  $T$  productions have a single terminal on the right. We use the constructor *Leaf* to create a suitable node which becomes the value of  $T\ node$ .

Figure 5.11 shows the construction of a syntax tree for the input  $a - 4 * c$ . The nodes of the syntax tree are shown as records with the *op* field. Syntax tree edges are now shown as solid lines. The underlying parse tree which need not actually be constructed is shown with dotted edges. The third type of line shown dashed represents the values of  $E\ node$  and  $T\ node$  each line points to the appropriate syntax tree node.

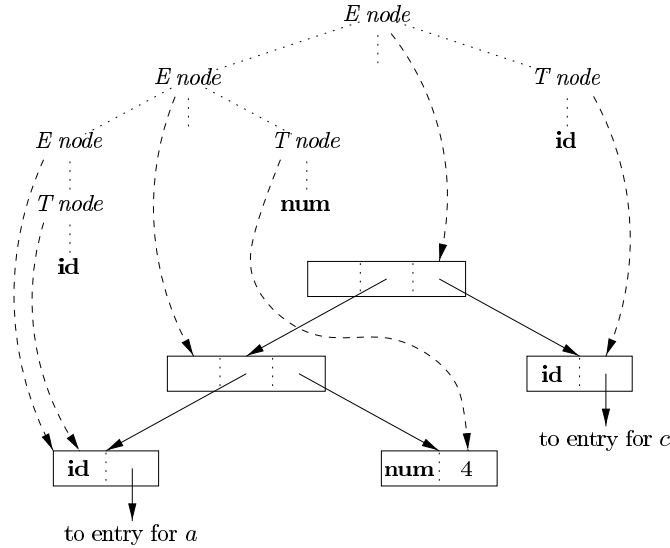
At the bottom we see leaves for  $a - 4$  and  $c$  constructed by *Leaf*. We suppose that the lexical value **id entry** points into the symbol table and the lexical value **num val** is the numerical value of a constant. These leaves or pointers to them become the value of  $T\ node$  at the three parse tree nodes labeled  $T$  according to rules 5 and 6. Note that by rule 3 the pointer to the leaf for  $a$  is also the value of  $E\ node$  for the leftmost  $E$  in the parse tree.

Rule 2 causes us to create a node with *op* equal to the minus sign and pointers to the first two leaves. Then rule 1 produces the root node of the syntax tree by combining the node for  $-$  with the third leaf.

If the rules are evaluated during a postorder traversal of the parse tree or with reductions during a bottom up parse then the sequence of steps shown in Fig. 5.12 ends with  $p_5$  pointing to the root of the constructed syntax tree.  $\square$

With a grammar designed for top down parsing the same syntax trees are constructed using the same sequence of steps even though the structure of the parse trees differs significantly from that of syntax trees.

**Example 5.12** The L attributed definition in Fig. 5.13 performs the same translation as the S attributed definition in Fig. 5.10. The attributes for the grammar symbols  $E$ ,  $T$ , **id** and **num** are as discussed in Example 5.11.


 Figure 5.11 Syntax tree for  $a * 4 + c$ 

```

1  p1  new Leaf id entry a
2  p2  new Leaf num 4
3  p3  new Node ' ' p1 p2
4  p4  new Leaf id entry c
5  p5  new Node ' ' p3 p4
    
```

 Figure 5.12 Steps in the construction of the syntax tree for  $a * 4 + c$ 

The rules for building syntax trees in this example are similar to the rules for the desk calculator in Example 5.3. In the desk calculator example a term  $x + y$  was evaluated by passing  $x$  as an inherited attribute since  $x$  and  $y$  appeared in different portions of the parse tree. Here the idea is to build a syntax tree for  $x + y$  by passing  $x$  as an inherited attribute since  $x$  and  $y$  appear in different subtrees. Nonterminal  $E'$  is the counterpart of nonterminal  $T'$  in Example 5.3. Compare the dependency graph for  $a * 4 + c$  in Fig. 5.14 with that for  $3 * 5$  in Fig. 5.7.

Nonterminal  $E'$  has an inherited attribute *inh* and a synthesized attribute *syn*. Attribute  $E' inh$  represents the partial syntax tree constructed so far. Specifically it represents the root of the tree for the prefix of the input string that is to the left of the subtree for  $E'$ . At node 5 in the dependency graph in Fig. 5.14  $E' inh$  denotes the root of the partial syntax tree for the identifier  $a$  that is the leaf for  $a$ . At node 6  $E' inh$  denotes the root for the partial syntax

### 5.3 APPLICATIONS OF SYNTAX DIRECTED TRANSLATION

PRODUCTION			SEMANTIC RULES
1	$E \rightarrow T E'$		$E \text{ node} \quad E' \text{ syn}$ $E' \text{ inh} \quad T \text{ node}$
2	$E' \rightarrow T E'_1$		$E'_1 \text{ inh} \quad \mathbf{new} \text{ Node } ' \quad ' E' \text{ inh } T \text{ node}$ $E' \text{ syn} \quad E'_1 \text{ syn}$
3	$E' \rightarrow T E'_1$		$E'_1 \text{ inh} \quad \mathbf{new} \text{ Node } ' \quad ' E' \text{ inh } T \text{ node}$ $E' \text{ syn} \quad E'_1 \text{ syn}$
4	$E' \rightarrow$		$E' \text{ syn} \quad E' \text{ inh}$
5	$T \rightarrow E$		$T \text{ node} \quad E \text{ node}$
6	$T \rightarrow \mathbf{id}$		$T \text{ node} \quad \mathbf{new} \text{ Leaf } \mathbf{id} \text{ id entry}$
7	$T \rightarrow \mathbf{num}$		$T \text{ node} \quad \mathbf{new} \text{ Leaf } \mathbf{num} \text{ num val}$

Figure 5.13 Constructing syntax trees during top down parsing




Figure 5.14 Dependency graph for  $a \ 4 \ c$  with the SDD of Fig. 5.13

tree for the input  $a \ 4 \ c$ . At node 9  $E' \text{ inh}$  denotes the syntax tree for  $a \ 4 \ c$ .

Since there is no more input at node 9  $E' \text{ inh}$  points to the root of the entire syntax tree. The *syn* attributes pass this value back up the parse tree until it becomes the value of  $E \text{ node}$ . Specifically the attribute value at node 10 is defined by the rule  $E' \text{ syn} \rightarrow E' \text{ inh}$  associated with the production  $E'$ . The attribute value at node 11 is defined by the rule  $E' \text{ syn} \rightarrow E'_1 \text{ syn}$  associated with production 2 in Fig. 5.13. Similar rules define the attribute values at nodes 12 and 13.  $\square$

#### 5.3.2 The Structure of a Type

Inherited attributes are useful when the structure of the parse tree differs from the abstract syntax of the input. Attributes can then be used to carry informa-

tion from one part of the parse tree to another. The next example shows how a mismatch in structure can be due to the design of the language and not due to constraints imposed by the parsing method.

**Example 5.13** In C the type `int 2 3` can be read as “array of 2 arrays of 3 integers”. The corresponding type expression `array 2 array 3 integer` is represented by the tree in Fig. 5.15. The operator `array` takes two parameters: a number and a type. If types are represented by trees, then this operator returns a tree node labeled `array` with two children: for a number and a type.




Figure 5.15 Type expression for `int 2 3`

With the SDD in Fig. 5.16, nonterminal  $T$  generates either a basic type or an array type. Nonterminal  $B$  generates one of the basic types `int` and `oat`.  $T$  generates a basic type when  $T$  derives  $BC$  and  $C$  derives  $\epsilon$ . Otherwise  $C$  generates array components consisting of a sequence of integers, each integer surrounded by brackets.

PRODUCTION		SEMANTIC RULES	
$T$	$BC$	$Tt \leftarrow Ct$	
		$Cb \leftarrow Bt$	
$B$	<code>int</code>	$Bt \leftarrow \text{integer}$	
$B$	<code>oat</code>	$Bt \leftarrow \text{oat}$	
$C$	<code>num</code> $C_1$	$Ct \leftarrow \text{array num val } C_1t$	
		$C_1b \leftarrow Cb$	
$C$		$Ct \leftarrow Cb$	

Figure 5.16  $T$  generates either a basic type or an array type

The nonterminals  $B$  and  $T$  have a synthesized attribute  $t$  representing a type. The nonterminal  $C$  has two attributes: an inherited attribute  $b$  and a synthesized attribute  $t$ . The inherited  $b$  attributes pass a basic type down the tree, and the synthesized  $t$  attributes accumulate the result.

An annotated parse tree for the input string `int 2 3` is shown in Fig. 5.17. The corresponding type expression in Fig. 5.15 is constructed by passing the type `integer` from  $B$  down the chain of  $C$ ’s through the inherited attributes  $b$ . The array type is synthesized up the chain of  $C$ ’s through the attributes  $t$ .

In more detail: at the root for  $T \rightarrow BC$ , nonterminal  $C$  inherits the type from  $B$  using the inherited attribute  $Cb$ . At the rightmost node for  $C$ , the



## 5.4 Syntax Directed Translation Schemes

Syntax directed translation schemes are a complementary notation to syntax directed definitions. All of the applications of syntax directed definitions in Section 5.3 can be implemented using syntax directed translation schemes.

From Section 2.3.5 a *syntax directed translation scheme* (SDT) is a context free grammar with program fragments embedded within production bodies. The program fragments are called *semantic actions* and can appear at any position within a production body. By convention, we place curly braces around actions if braces are needed as grammar symbols; then we quote them.

Any SDT can be implemented by first building a parse tree and then performing the actions in a left to right depth first order—that is, during a preorder traversal. An example appears in Section 5.4.3.

Typically, SDTs are implemented during parsing, without building a parse tree. In this section, we focus on the use of SDTs to implement two important classes of SDDs:

1. The underlying grammar is LR parsable, and the SDD is S attributed.
2. The underlying grammar is LL parsable, and the SDD is L attributed.

We shall see how, in both these cases, the semantic rules in an SDD can be converted into an SDT with actions that are executed at the right time. During parsing, an action in a production body is executed as soon as all the grammar symbols to the left of the action have been matched.

SDTs that can be implemented during parsing can be characterized by introducing distinct *marker nonterminals* in place of each embedded action: each marker  $M$  has only one production,  $M \rightarrow \dots$ . If the grammar with marker nonterminals can be parsed by a given method, then the SDT can be implemented during parsing.

### 5.4.1 Postfix Translation Schemes

By far the simplest SDD implementation occurs when we can parse the grammar bottom up and the SDD is S attributed. In that case, we can construct an SDT in which each action is placed at the end of the production and is executed along with the reduction of the body to the head of that production. SDTs with all actions at the right ends of the production bodies are called *postfix SDTs*.

**Example 5.14** The postfix SDT in Fig. 5.18 implements the desk calculator SDD of Fig. 5.1, with one change: the action for the first production prints a value. The remaining actions are exact counterparts of the semantic rules. Since the underlying grammar is LR, and the SDD is S attributed, these actions can be correctly performed along with the reduction steps of the parser.  $\square$

## 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

$L$	$E\ \mathbf{n}$	$\{ \text{print } E\ val \}$
$E$	$E_1\ T$	$\{ E\ val\ E_1\ val\ T\ val \}$
$E$	$T$	$\{ E\ val\ T\ val \}$
$T$	$T_1\ F$	$\{ T\ val\ T_1\ val\ F\ val \}$
$T$	$F$	$\{ T\ val\ F\ val \}$
$F$	$E$	$\{ F\ val\ E\ val \}$
$F$	$\mathbf{digit}$	$\{ F\ val\ \mathbf{digit}\ lexval \}$

Figure 5.18 Postfix SDT implementing the desk calculator

### 5.4.2 Parser Stack Implementation of Postfix SDTs

Postfix SDTs can be implemented during LR parsing by executing the actions when reductions occur. The attributes of each grammar symbol can be put on the stack in a place where they can be found during the reduction. The best plan is to place the attributes along with the grammar symbols or the LR states that represent these symbols in records on the stack itself.

In Fig. 5.19 the parser stack contains records with a field for a grammar symbol or parser state and below it a field for an attribute. The three grammar symbols  $XYZ$  are on top of the stack; perhaps they are about to be reduced according to a production like  $A \rightarrow XYZ$ . Here we show  $Xx$  as the one attribute of  $X$  and so on. In general we can allow for more attributes either by making the records large enough or by putting pointers to records on the stack. With small attributes it may be simpler to make the records large enough even if some fields go unused some of the time. However, if one or more attributes are of unbounded size (say they are character strings) then it would be better to put a pointer to the attribute's value in the stack record and store the actual value in some larger shared storage area that is not part of the stack.

	$X$	$Y$	$Z$	State: grammar symbol
	$Xx$	$Yy$	$Zz$	Synthesized attributes
			↑	
			top	

Figure 5.19 Parser stack with a field for synthesized attributes

If the attributes are all synthesized and the actions occur at the ends of the productions then we can compute the attributes for the head when we reduce the body to the head. If we reduce by a production such as  $A \rightarrow XYZ$  then we have all the attributes of  $X$ ,  $Y$ , and  $Z$  available at known positions on the stack as in Fig. 5.19. After the action  $A$  and its attributes are at the top of the stack in the position of the record for  $X$ .

**Example 5.15** Let us rewrite the actions of the desk calculator SDT of Ex

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

ample 5.14 so that they manipulate the parser stack explicitly. Such stack manipulation is usually done automatically by the parser.

PRODUCTION	ACTIONS
$L \rightarrow E \mathbf{n}$	{ print $stack[top - 1].val$ $top = top - 1$ }
$E \rightarrow E_1 T$	{ $stack[top - 2].val = stack[top - 2].val \text{ } stack[top].val$ $top = top - 2$ }
$E \rightarrow T$	
$T \rightarrow T_1 F$	{ $stack[top - 2].val = stack[top - 2].val \text{ } stack[top].val$ $top = top - 2$ }
$T \rightarrow F$	
$F \rightarrow E$	{ $stack[top - 2].val = stack[top - 1].val$ $top = top - 2$ }
$F \rightarrow \mathbf{digit}$	

Figure 5.20 Implementing the desk calculator on a bottom up parsing stack

Suppose that the stack is kept in an array of records called *stack* with *top* a cursor to the top of the stack. Thus *stack[top]* refers to the top record on the stack, *stack[top - 1]* to the record below that, and so on. Also, we assume that each record has a field called *val* which holds the attribute of whatever grammar symbol is represented in that record. Thus, we may refer to the attribute *E.val* that appears at the third position on the stack as *stack[top - 2].val*. The entire SDT is shown in Fig. 5.20.

For instance, in the second production  $E \rightarrow E_1 T$  we go two positions below the top to get the value of  $E_1$  and we find the value of  $T$  at the top. The resulting sum is placed where the head  $E$  will appear after the reduction—that is, two positions below the current top. The reason is that after the reduction the three topmost stack symbols are replaced by one. After computing  $E.val$  we pop two symbols off the top of the stack, so the record where we placed  $E.val$  will now be at the top of the stack.

In the third production  $E \rightarrow T$  no action is necessary, because the length of the stack does not change, and the value of  $T.val$  at the stack top will simply become the value of  $E.val$ . The same observation applies to the productions  $T \rightarrow F$  and  $F \rightarrow \mathbf{digit}$ . Production  $F \rightarrow E$  is slightly different. Although the value does not change, two positions are removed from the stack during the reduction, so the value has to move to the position after the reduction.

Note that we have omitted the steps that manipulate the first field of the stack records—the field that gives the LR state or otherwise represents the grammar symbol. If we are performing an LR parse, the parsing table tells us what the new state is every time we reduce (see Algorithm 4.44). Thus, we may



## 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

simply place that state in the record for the new top of stack  $\square$

### 5.4.3 SDT's With Actions Inside Productions

An action may be placed at any position within the body of a production. It is performed immediately after all symbols to its left are processed. Thus if we have a production  $B \rightarrow X \{a\} Y$  the action  $a$  is done after we have recognized  $X$  if  $X$  is a terminal or all the terminals derived from  $X$  if  $X$  is a nonterminal. More precisely

If the parse is bottom up then we perform action  $a$  as soon as this occurrence of  $X$  appears on the top of the parsing stack

If the parse is top down we perform  $a$  just before we attempt to expand this occurrence of  $Y$  if  $Y$  a nonterminal or check for  $Y$  on the input if  $Y$  is a terminal

SDT's that can be implemented during parsing include post-fix SDT's and a class of SDT's considered in Section 5.5 that implements L attributed definitions. Not all SDT's can be implemented during parsing as we shall see in the next example

**Example 5.16** As an extreme example of a problematic SDT suppose that we turn our desk calculator running example into an SDT that prints the prefix form of an expression rather than evaluating the expression. The productions and actions are shown in Fig. 5.21

1	$L$	$E \mathbf{n}$
2	$E$	$\{ \text{print ' ' } \} E_1 T$
3	$E$	$T$
4	$T$	$\{ \text{print ' ' } \} T_1 F$
5	$T$	$F$
6	$F$	$E$
7	$F$	$\mathbf{digit} \{ \text{print } \mathbf{digit} \text{ lexical } \}$

Figure 5.21 Problematic SDT for infix to prefix translation during parsing

Unfortunately it is impossible to implement this SDT during either top down or bottom up parsing because the parser would have to perform critical actions like printing instances of  $\mathbf{digit}$  or  $\mathbf{lexical}$  long before it knows whether these symbols will appear in its input

Using marker nonterminals  $M_2$  and  $M_4$  for the actions in productions 2 and 4 respectively on input that is a digit a shift reduce parser see Section 4.5.3 has conflicts between reducing by  $M_2$  reducing by  $M_4$  and shifting the digit  $\square$

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

Any SDT can be implemented as follows

- 1 Ignoring the actions parse the input and produce a parse tree as a result
- 2 Then examine each interior node  $N$  say one for production  $A$  Add additional children to  $N$  for the actions in so the children of  $N$  from left to right have exactly the symbols and actions of
- 3 Perform a preorder traversal see Section 2.3.4 of the tree and as soon as a node labeled by an action is visited perform that action

For instance Fig. 5.22 shows the parse tree for expression  $3 \cdot 5 \cdot 4$  with actions inserted. If we visit the nodes in preorder we get the prefix form of the expression  $\cdot 3 \cdot 5 \cdot 4$

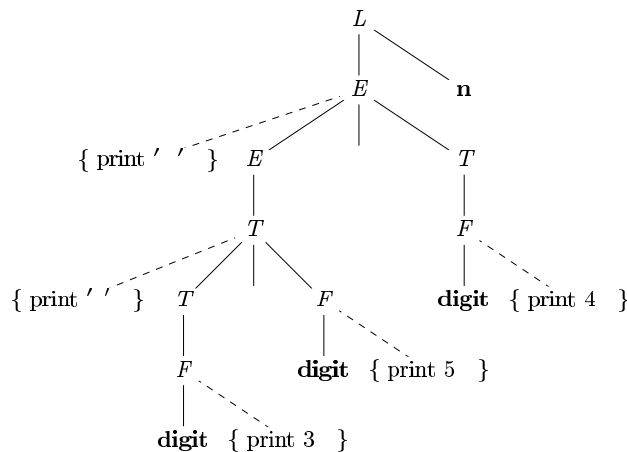


Figure 5.22 Parse tree with actions embedded

### 5.4.4 Eliminating Left Recursion From SDT's

Since no grammar with left recursion can be parsed deterministically top down we examined left recursion elimination in Section 4.3.3. When the grammar is part of an SDT we also need to worry about how the actions are handled.

First consider the simple case in which the only thing we care about is the order in which the actions in an SDT are performed. For example if each action simply prints a string we care only about the order in which the strings are printed. In this case the following principle can guide us.

When transforming the grammar treat the actions as if they were terminal symbols.

#### 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

This principle is based on the idea that the grammar transformation preserves the order of the terminals in the generated string. The actions are therefore executed in the same order in any left to right parse: top down or bottom up.

The "trick" for eliminating left recursion is to take two productions

$$A \rightarrow A \mid$$

that generate strings consisting of a and any number of s and replace them by productions that generate the same strings using a new nonterminal  $R$  for remainder of the first production

$$\begin{array}{l} A \rightarrow R \\ R \rightarrow R \mid \end{array}$$

If does not begin with  $A$  then  $A$  no longer has a left recursive production. In regular definition terms, with both sets of productions,  $A$  is defined by. See Section 4.3.3 for the handling of situations where  $A$  has more recursive or nonrecursive productions.

**Example 5.17** Consider the following  $E$  productions from an SDT for translating infix expressions into postfix notation

$$\begin{array}{l} E \rightarrow E_1 T \{ \text{print ' ' } \} \\ E \rightarrow T \end{array}$$

If we apply the standard transformation to  $E$ , the remainder of the left recursive production is

$$T \{ \text{print ' ' } \}$$

and the body of the other production is  $T$ . If we introduce  $R$  for the remainder of  $E$ , we get the set of productions

$$\begin{array}{l} E \rightarrow T R \\ R \rightarrow T \{ \text{print ' ' } \} R \\ R \end{array}$$

□

When the actions of an SDD compute attributes rather than merely printing output, we must be more careful about how we eliminate left recursion from a grammar. However, if the SDD is S-attributed, then we can always construct an SDT by placing attribute computing actions at appropriate positions in the new productions.

We shall give a general schema for the case of a single recursive production, a single nonrecursive production, and a single attribute of the left recursive nonterminal. The generalization to many productions of each type is not hard but is notationally cumbersome. Suppose that the two productions are

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

$$\begin{array}{l} A \quad A_1 Y \{A a \quad g A_1 a Y y\} \\ A \quad X \{A a \quad f X x\} \end{array}$$

Here  $A a$  is the synthesized attribute of left recursive nonterminal  $A$  and  $X$  and  $Y$  are single grammar symbols with synthesized attributes  $X x$  and  $Y y$  respectively. These could represent a string of several grammar symbols, each with its own attributes, since the schema has an arbitrary function  $g$  computing  $A a$  in the recursive production and an arbitrary function  $f$  computing  $A a$  in the second production. In each case,  $f$  and  $g$  take as arguments whatever attributes they are allowed to access if the SDD is S attributed.

We want to turn the underlying grammar into

$$\begin{array}{l} A \quad X R \\ R \quad Y R \mid \end{array}$$

Figure 5.23 suggests what the SDT on the new grammar must do. In (a) we see the effect of the post-fix SDT on the original grammar. We apply  $f$  once corresponding to the use of production  $A \rightarrow X$  and then apply  $g$  as many times as we use the production  $A \rightarrow A Y$ . Since  $R$  generates a remainder of  $Y$ 's, its translation depends on the string to its left, a string of the form  $XY Y^* Y$ . Each use of the production  $R \rightarrow Y R$  results in an application of  $g$ . For  $R$  we use an inherited attribute  $R i$  to accumulate the result of successively applying  $g$ , starting with the value of  $A a$ .




Figure 5.23 Eliminating left recursion from a post-fix SDT

In addition,  $R$  has a synthesized attribute  $R s$  not shown in Fig. 5.23. This attribute is first computed when  $R$  ends its generation of  $Y$  symbols, as signaled by the use of production  $R \rightarrow Y R$ .  $R s$  is then copied up the tree, so it can become the value of  $A a$  for the entire expression  $XY Y^* Y$ . The case where  $A$  generates  $XY Y^* Y$  is shown in Fig. 5.23, and we see that the value of  $A a$  at the root of (a) has two uses of  $g$ . So does  $R i$  at the bottom of tree (b), and it is this value of  $R s$  that gets copied up that tree.

To accomplish this translation, we use the following SDT

## 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

$$\begin{array}{ll}
 A & X \{R i \quad f X x \} R \{A a \quad R s\} \\
 R & Y \{R_1 i \quad g R i Y y \} R_1 \{R s \quad R_1 s\} \\
 R & \{R s \quad R i\}
 \end{array}$$

Notice that the inherited attribute  $R i$  is evaluated immediately before a use of  $R$  in the body while the synthesized attributes  $A a$  and  $R s$  are evaluated at the ends of the productions. Thus whatever values are needed to compute these attributes will be available from what has been computed to the left.

### 5.4.5 SDT's for L Attributed Definitions

In Section 5.4.1 we converted S attributed SDD's into post-x SDT's with actions at the right ends of productions. As long as the underlying grammar is LR, post-x SDT's can be parsed and translated bottom up.

Now we consider the more general case of an L attributed SDD. We shall assume that the underlying grammar can be parsed top down, for if not it is frequently impossible to perform the translation in connection with either an LL or an LR parser. With any grammar the technique below can be implemented by attaching actions to a parse tree and executing them during preorder traversal of the tree.

The rules for turning an L attributed SDD into an SDT are as follows:

1. Embed the action that computes the inherited attributes for a nonterminal  $A$  immediately before that occurrence of  $A$  in the body of the production. If several inherited attributes for  $A$  depend on one another in an acyclic fashion, order the evaluation of attributes so that those needed first are computed first.
2. Place the actions that compute a synthesized attribute for the head of a production at the end of the body of that production.

We shall illustrate these principles with two extended examples. The first involves typesetting. It illustrates how the techniques of compiling can be used in language processing for applications other than what we normally think of as programming languages. The second example is about the generation of intermediate code for a typical programming language construct—a form of while statement.

**Example 5.18** This example is motivated by languages for typesetting mathematical formulas. *Eqn* is an early example of such a language; ideas from *Eqn* are still found in the *T<sub>E</sub>X* typesetting system, which was used to produce this book.

We shall concentrate on only the capability to define subscripts, subscripts of subscripts, and so on, ignoring superscripts, built up fractions, and all other mathematical features. In the *Eqn* language, one writes `a sub i sub j` to set the expression  $a_{ij}$ . A simple grammar for *boxes*—elements of text bounded by a rectangle—is

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

$$B \rightarrow B_1 B_2 \mid B_1 \mathbf{sub} B_2 \mid (B_1) \mid \mathbf{text}$$

Corresponding to these four productions a box can be either

- 1 Two boxes juxtaposed with the first  $B_1$  to the left of the other  $B_2$
- 2 A box and a subscript box. The second box appears in a smaller size lower and to the right of the first box
- 3 A parenthesized box for grouping of boxes and subscripts. Eqn and T<sub>E</sub>X both use curly braces for grouping but we shall use ordinary round parentheses to avoid confusion with the braces that surround actions in SDT's
- 4 A text string that is any string of characters

This grammar is ambiguous but we can still use it to parse bottom up if we make subscripting and juxtaposition right associative with **sub** taking precedence over juxtaposition

Expressions will be typeset by constructing larger boxes out of smaller ones. In Fig. 5.24 the boxes for  $E_1$  and *height* are about to be juxtaposed to form the box for  $E_1 \text{ height}$ . The left box for  $E_1$  is itself constructed from the box for  $E$  and the subscript 1. The subscript 1 is handled by shrinking its box by about 30% lowering it and placing it after the box for  $E$ . Although we shall treat *height* as a text string the rectangles within its box show how it can be constructed from boxes for the individual letters.




Figure 5.24 Constructing larger boxes from smaller ones

In this example we concentrate on the vertical geometry of boxes only. The horizontal geometry—the widths of boxes—is also interesting, especially when different characters have different widths. It may not be readily apparent but each of the distinct characters in Fig. 5.24 has a different width.

The values associated with the vertical geometry of boxes are as follows

- a The *point size* is used to set text within a box. We shall assume that characters not in subscripts are set in 10 point type—the size of type in this book. Further we assume that if a box has point size  $p$  then its subscript box has the smaller point size  $0.7p$ . Inherited attribute  $B.ps$  will represent the point size of block  $B$ . This attribute must be inherited because the context determines by how much a given box needs to be shrunk due to the number of levels of subscripting.

#### 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

- b Each box has a *baseline* which is a vertical position that corresponds to the bottoms of lines of text not counting any letters like *g* that extend below the normal baseline In Fig 5.24 the dotted line represents the baseline for the boxes *E* *height* and the entire expression The baseline for the box containing the subscript 1 is adjusted to lower the subscript
- c A box has a *height* which is the distance from the top of the box to the baseline Synthesized attribute *B ht* gives the height of box *B*
- d A box has a *depth* which is the distance from the baseline to the bottom of the box Synthesized attribute *B dp* gives the depth of box *B*

The SDD in Fig 5.25 gives rules for computing point sizes heights and depths Production 1 is used to assign *B ps* the initial value 10

PRODUCTION	SEMANTIC RULES
1 <i>S B</i>	<i>B ps</i> 10
2 <i>B B<sub>1</sub> B<sub>2</sub></i>	<i>B<sub>1</sub> ps B ps</i> <i>B<sub>2</sub> ps B ps</i> <i>B ht max B<sub>1</sub> ht B<sub>2</sub> ht</i> <i>B dp max B<sub>1</sub> dp B<sub>2</sub> dp</i>
3 <i>B B<sub>1</sub> sub B<sub>2</sub></i>	<i>B<sub>1</sub> ps B ps</i> <i>B<sub>2</sub> ps 0.7 B ps</i> <i>B ht max B<sub>1</sub> ht B<sub>2</sub> ht 0.25 B ps</i> <i>B dp max B<sub>1</sub> dp B<sub>2</sub> dp 0.25 B ps</i>
4 <i>B B<sub>1</sub></i>	<i>B<sub>1</sub> ps B ps</i> <i>B ht B<sub>1</sub> ht</i> <i>B dp B<sub>1</sub> dp</i>
5 <i>B text</i>	<i>B ht getHt B ps text lerval</i> <i>B dp getDp B ps text lerval</i>

Figure 5.25 SDD for typesetting boxes

Production 2 handles juxtaposition Point sizes are copied down the parse tree that is two sub boxes of a box inherit the same point size from the larger box Heights and depths are computed up the tree by taking the maximum That is the height of the larger box is the maximum of the heights of its two components and similarly for the depth

Production 3 handles subscripting and is the most subtle In this greatly simplified example we assume that the point size of a subscripted box is 70% of the point size of its parent Reality is much more complex since subscripts cannot shrink indefinitely in practice after a few levels the sizes of subscripts

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

shrink hardly at all. Further, we assume that the baseline of a subscript box drops by 25% of the parent's point size; again, reality is more complex.

Production 4 copies attributes appropriately when parentheses are used. Finally, production 5 handles the leaves that represent text boxes. In this matter too, the true situation is complicated, so we merely show two unspecified functions *getHt* and *getDp* that examine tables created with each font to determine the maximum height and maximum depth of any characters in the text string. The string itself is presumed to be provided as the attribute *lexval* of terminal **text**.

Our last task is to turn this SDD into an SDT, following the rules for an L-attributed SDD, which Fig. 5.25 is. The appropriate SDT is shown in Fig. 5.26. For readability, since production bodies become long, we split them across lines and line up the actions. Production bodies therefore consist of the contents of all lines up to the head of the next production.  $\square$

	PRODUCTION	ACTIONS
1	<i>S</i>	{ <i>B ps</i> 10 }
	<i>B</i>	
2	<i>B</i>	{ <i>B<sub>1</sub> ps</i> <i>B ps</i> }
	<i>B<sub>1</sub></i>	{ <i>B<sub>2</sub> ps</i> <i>B ps</i> }
	<i>B<sub>2</sub></i>	{ <i>B ht</i> max <i>B<sub>1</sub> ht</i> <i>B<sub>2</sub> ht</i> <i>B dp</i> max <i>B<sub>1</sub> dp</i> <i>B<sub>2</sub> dp</i> }
3	<i>B</i>	{ <i>B<sub>1</sub> ps</i> <i>B ps</i> }
	<i>B<sub>1</sub> sub</i>	{ <i>B<sub>2</sub> ps</i> 0.7 <i>B ps</i> }
	<i>B<sub>2</sub></i>	{ <i>B ht</i> max <i>B<sub>1</sub> ht</i> <i>B<sub>2</sub> ht</i> 0.25 <i>B ps</i> <i>B dp</i> max <i>B<sub>1</sub> dp</i> <i>B<sub>2</sub> dp</i> 0.25 <i>B ps</i> }
4	<i>B</i>	{ <i>B<sub>1</sub> ps</i> <i>B ps</i> }
	<i>B<sub>1</sub></i>	{ <i>B ht</i> <i>B<sub>1</sub> ht</i> <i>B dp</i> <i>B<sub>1</sub> dp</i> }
5	<i>B</i> <b>text</b>	{ <i>B ht</i> <i>getHt</i> <i>B ps</i> <b>text</b> <i>lexval</i> <i>B dp</i> <i>getDp</i> <i>B ps</i> <b>text</b> <i>lexval</i> }

Figure 5.26 SDT for typesetting boxes

Our next example concentrates on a simple while statement and the generation of intermediate code for this type of statement. Intermediate code will be treated as a string-valued attribute. Later, we shall explore techniques that involve the writing of pieces of a string-valued attribute as we parse, thus avoiding the copying of long strings to build even longer strings. The technique was introduced in Example 5.17, where we generated the postfix form of an infix



## 5.4 SYNTAX DIRECTED TRANSLATION SCHEMES

expression on the fly rather than computing it as an attribute. However, in our first formulation, we create a string valued attribute by concatenation.

**Example 5.19** For this example, we only need one production.

$$S \rightarrow \text{while } C \text{ } S_1$$

Here  $S$  is the nonterminal that generates all kinds of statements, presumably including if statements, assignment statements, and others. In this example,  $C$  stands for a conditional expression, a boolean expression that evaluates to true or false.

In this flow of control example, the only things we ever generate are labels. All the other intermediate code instructions are assumed to be generated by parts of the SDT that are not shown. Specifically, we generate explicit instructions of the form **label**  $L$ , where  $L$  is an identifier to indicate that  $L$  is the label of the instruction that follows. We assume that the intermediate code is like that introduced in Section 2.8.4.

The meaning of our while statement is that the conditional  $C$  is evaluated. If it is true, control goes to the beginning of the code for  $S_1$ . If false, then control goes to the code that follows the while statement's code. The code for  $S_1$  must be designed to jump to the beginning of the code for the while statement when finished; the jump to the beginning of the code that evaluates  $C$  is not shown in Fig. 5.27.

We use the following attributes to generate the proper intermediate code.

1. The inherited attribute  $S_{next}$  labels the beginning of the code that must be executed after  $S$  is finished.
2. The synthesized attribute  $S_{code}$  is the sequence of intermediate code steps that implements a statement  $S$  and ends with a jump to  $S_{next}$ .
3. The inherited attribute  $C_{true}$  labels the beginning of the code that must be executed if  $C$  is true.
4. The inherited attribute  $C_{false}$  labels the beginning of the code that must be executed if  $C$  is false.
5. The synthesized attribute  $C_{code}$  is the sequence of intermediate code steps that implements the condition  $C$  and jumps either to  $C_{true}$  or to  $C_{false}$  depending on whether  $C$  is true or false.

The SDD that computes these attributes for the while statement is shown in Fig. 5.27. A number of points merit explanation.

The function *new* generates new labels.

The variables  $L1$  and  $L2$  hold labels that we need in the code.  $L1$  is the beginning of the code for the while statement, and we need to arrange

## CHAPTER 5 SYNTAX DIRECTED TRANSLATION

```

S  while C S1  L1  new
                      L2  new
                      S1 next  L1
                      C false  S next
                      C true   L2
                      S code   label || L1 || C code || label || L2 || S1 code

```

Figure 5 27 SDD for while statements

that  $S_1$  jumps there after it finishes. That is why we set  $S_1 \text{ next}$  to  $L1$ .  $L2$  is the beginning of the code for  $S_1$  and it becomes the value of  $C \text{ true}$  because we branch there when  $C$  is true.

Notice that  $C \text{ false}$  is set to  $S \text{ next}$  because when the condition is false we execute whatever code must follow the code for  $S$ .

We use  $\parallel$  as the symbol for concatenation of intermediate code fragments. The value of  $S \text{ code}$  thus begins with the label  $L1$ , then the code for condition  $C$ , another label  $L2$ , and the code for  $S_1$ .

This SDD is L attributed. When we convert it into an SDT, the only remaining issue is how to handle the labels  $L1$  and  $L2$ , which are variables and not attributes. If we treat actions as dummy nonterminals, then such variables can be treated as the synthesized attributes of dummy nonterminals. Since  $L1$  and  $L2$  do not depend on any other attributes, they can be assigned to the first action in the production. The resulting SDT with embedded actions that implements this L attributed definition is shown in Fig. 5 28.  $\square$

```

S  while  { L1  new  L2  new  C false  S next  C true  L2  }
C        { S1 next  L1  }
S1      { S code   label || L1 || C code || label || L2 || S1 code  }

```

Figure 5 28 SDT for while statements

### 5 4 6 Exercises for Section 5 4

**Exercise 5 4 1** We mentioned in Section 5 4 2 that it is possible to deduce from the LR state on the parsing stack what grammar symbol is represented by the state. How would we discover this information?

**Exercise 5 4 2** Rewrite the following SDT

$$\begin{array}{lcl}
 A & A \{a\} B & | A B \{b\} | 0 \\
 B & B \{c\} A & | B A \{d\} | 1
 \end{array}$$

## 5.5 IMPLEMENTING L ATTRIBUTED SDD S

so that the underlying grammar becomes non left recursive. Here  $a$ ,  $b$ ,  $c$  and  $d$  are actions and 0 and 1 are terminals.

**Exercise 5.4.3** The following SDT computes the value of a string of 0's and 1's interpreted as a positive binary integer.

$$\begin{array}{lcl}
 B & \rightarrow & B_1 0 \{B.val \cdot 2, B_1.val\} \\
 & | & B_1 1 \{B.val \cdot 2, B_1.val + 1\} \\
 & | & 1 \{B.val + 1\}
 \end{array}$$

Rewrite this SDT so the underlying grammar is not left recursive and yet the same value of  $B.val$  is computed for the entire input string.

**Exercise 5.4.4** Write L attributed SDD's analogous to that of Example 5.19 for the following productions, each of which represents a familiar flow of control construct as in the programming language C. You may need to generate a three address statement to jump to a particular label  $L$  in which case you should generate **goto**  $L$ .

- a.  $S \rightarrow \text{if } C \text{ } S_1 \text{ else } S_2$
- b.  $S \rightarrow \text{do } S_1 \text{ while } C$
- c.  $S \rightarrow \{ 'L' \}^* L \mid L S \mid$

Note that any statement in the list can have a jump from its middle to the next statement, so it is not sufficient simply to generate code for each statement in order.

**Exercise 5.4.5** Convert each of your SDD's from Exercise 5.4.4 to an SDT in the manner of Example 5.19.

**Exercise 5.4.6** Modify the SDD of Fig. 5.25 to include a synthesized attribute  $Box.le$  the length of a box. The length of the concatenation of two boxes is the sum of the lengths of each. Then add your new rules to the proper positions in the SDT of Fig. 5.26.

**Exercise 5.4.7** Modify the SDD of Fig. 5.25 to include superscripts denoted by operator **sup** between boxes. If box  $B_2$  is a superscript of box  $B_1$  then position the baseline of  $B_2$  0.6 times the point size of  $B_1$  above the baseline of  $B_1$ . Add the new production and rules to the SDT of Fig. 5.26.

## 5.5 Implementing L Attributed SDD's

Since many translation applications can be addressed using L attributed definitions, we shall consider their implementation in more detail in this section. The following methods do translation by traversing a parse tree.